

Computational aspects of finding Nash Equilibria for 2-player games

Maria Serna

Fall 2019

- 1 Linear Algebra formulation
- 2 Zero-sum games
- 3 The complexity of finding a NE
- 4 An exact algorithm to compute NE
- 5 Other algorithms

Nash equilibrium

Consider a 2-player game $\Gamma = (A_1, A_2, u_1, u_2)$.

Let $X = \Delta(A_1)$ and $Y = \Delta(A_2)$.

($\Delta(A)$ is the set of probability distributions over A)

Nash equilibrium

Consider a 2-player game $\Gamma = (A_1, A_2, u_1, u_2)$.

Let $X = \Delta(A_1)$ and $Y = \Delta(A_2)$.

($\Delta(A)$ is the set of probability distributions over A)

A **Nash equilibrium** is a mixed strategy profile $\sigma = (x, y) \in X \times Y$ such that, for every $x' \in X$, $y' \in Y$, it holds

$$U_1(x, y) \geq U_1(x', y) \text{ and } U_2(x, y) \geq U_2(x, y')$$

Linear algebra notation

Consider a 2-player game $\Gamma = (A_1, A_2, u_1, u_2)$

Linear algebra notation

Consider a 2-player game $\Gamma = (A_1, A_2, u_1, u_2)$

Let $n = |A_1|$ and $m = |A_2|$.

Linear algebra notation

Consider a 2-player game $\Gamma = (A_1, A_2, u_1, u_2)$

Let $n = |A_1|$ and $m = |A_2|$.

$x \in X$ is a **probability distribution**:

$$x = (x_1, \dots, x_n), x_i \geq 0 \text{ and } x_1 + \dots + x_n = 1$$

Linear algebra notation

Consider a 2-player game $\Gamma = (A_1, A_2, u_1, u_2)$

Let $n = |A_1|$ and $m = |A_2|$.

$x \in X$ is a **probability distribution**:

$$x = (x_1, \dots, x_n), x_i \geq 0 \text{ and } x_1 + \dots + x_n = 1$$

$y \in Y$ is a **probability distribution**:

$$y = (y_1, \dots, y_m), y_j \geq 0 \text{ and } y_1 + \dots + y_m = 1$$

Linear algebra notation

Consider a 2-player game $\Gamma = (A_1, A_2, u_1, u_2)$

Let $n = |A_1|$ and $m = |A_2|$.

$x \in X$ is a **probability distribution**:

$$x = (x_1, \dots, x_n), x_i \geq 0 \text{ and } x_1 + \dots + x_n = 1$$

$y \in Y$ is a **probability distribution**:

$$y = (y_1, \dots, y_m), y_j \geq 0 \text{ and } y_1 + \dots + y_m = 1$$

Utilities can be described by a $n \times m$ matrix R , for the row player, and C , for the column player. Then,

Linear algebra notation

Consider a 2-player game $\Gamma = (A_1, A_2, u_1, u_2)$

Let $n = |A_1|$ and $m = |A_2|$.

$x \in X$ is a **probability distribution**:

$$x = (x_1, \dots, x_n), x_i \geq 0 \text{ and } x_1 + \dots + x_n = 1$$

$y \in Y$ is a **probability distribution**:

$$y = (y_1, \dots, y_m), y_j \geq 0 \text{ and } y_1 + \dots + y_m = 1$$

Utilities can be described by a $n \times m$ matrix R , for the row player, and C , for the column player. Then,

$$U_1(x, y) = x^T R y \text{ and } U_2(x, y) = x^T C y$$

Computing a best response

For a given $x \in X$, we have to solve:

Computing a best response

For a given $x \in X$, we have to solve:

$$\begin{aligned} & \max \quad x^T R y \\ \text{Subject to: } & y_1 + \cdots + y_m = 1, y_j \geq 0. \end{aligned}$$

Computing a best response

For a given $x \in X$, we have to solve:

$$\begin{aligned} & \max \quad x^T R y \\ \text{Subject to: } & y_1 + \cdots + y_m = 1, y_j \geq 0. \end{aligned}$$

For a given y , we have to solve:

Computing a best response

For a given $x \in X$, we have to solve:

$$\begin{aligned} & \max \quad x^T R y \\ & \text{Subject to: } y_1 + \cdots + y_m = 1, y_j \geq 0. \end{aligned}$$

For a given y , we have to solve:

$$\begin{aligned} & \max \quad x^T C y \\ & \text{Subject to: } x_1 + \cdots + x_n = 1, x_i \geq 0 \end{aligned}$$

Computing a best response

For a given $x \in X$, we have to solve:

$$\begin{aligned} & \max \quad x^T R y \\ & \text{Subject to: } y_1 + \cdots + y_m = 1, y_j \geq 0. \end{aligned}$$

For a given y , we have to solve:

$$\begin{aligned} & \max \quad x^T C y \\ & \text{Subject to: } x_1 + \cdots + x_n = 1, x_i \geq 0 \end{aligned}$$

Those are linear programming problems, so

Computing a best response

For a given $x \in X$, we have to solve:

$$\begin{aligned} & \max \quad x^T R y \\ & \text{Subject to: } y_1 + \cdots + y_m = 1, y_j \geq 0. \end{aligned}$$

For a given y , we have to solve:

$$\begin{aligned} & \max \quad x^T C y \\ & \text{Subject to: } x_1 + \cdots + x_n = 1, x_i \geq 0 \end{aligned}$$

Those are linear programming problems, so

A best response can be computed in polynomial time for 2-player games with rational utilities.

- 1 Linear Algebra formulation
- 2 Zero-sum games**
- 3 The complexity of finding a NE
- 4 An exact algorithm to compute NE
- 5 Other algorithms

Zero-sum games

- A **zero-sum** game is a 2-player game such that, for each pure strategy profile (a, b) , $u_1(a, b) + u_2(a, b) = 0$.

Zero-sum games

- A **zero-sum** game is a 2-player game such that, for each pure strategy profile (a, b) , $u_1(a, b) + u_2(a, b) = 0$.
- That is **Let $u = u_1$, we have $u_2 = -u$.**

Zero-sum games

- A **zero-sum** game is a 2-player game such that, for each pure strategy profile (a, b) , $u_1(a, b) + u_2(a, b) = 0$.
- That is **Let $u = u_1$, we have $u_2 = -u$.**
- Player 1 is interested in **maximizing** u and player 2 in **minimizing** u .

Zero-sum games

- A **zero-sum** game is a 2-player game such that, for each pure strategy profile (a, b) , $u_1(a, b) + u_2(a, b) = 0$.
- That is **Let $u = u_1$, we have $u_2 = -u$.**
- Player 1 is interested in **maximizing u** and player 2 in **minimizing u .**
- In terms of matrices we have **$C = -R$.**

ZS: Nash conditions

- (x^*, y^*) is a NE

ZS: Nash conditions

- (x^*, y^*) is a NE

$$\begin{aligned}(x^*)^T R y^* &\geq x^T R y^*, \text{ for } x \in X, \\ (x^*)^T C y^* &\geq (x^*)^T C y, \text{ for } y \in Y.\end{aligned}$$

ZS: Nash conditions

- (x^*, y^*) is a NE

$$(x^*)^T R y^* \geq x^T R y^*, \text{ for } x \in X,$$

$$(x^*)^T C y^* \geq (x^*)^T C y, \text{ for } y \in Y.$$

- As $C = -R$ the second equation becomes

ZS: Nash conditions

- (x^*, y^*) is a NE
$$(x^*)^T R y^* \geq x^T R y^*, \text{ for } x \in X,$$
$$(x^*)^T C y^* \geq (x^*)^T C y, \text{ for } y \in Y.$$
- As $C = -R$ the second equation becomes
$$(x^*)^T R y^* \leq (x^*)^T R y, \text{ for } y \in Y.$$

ZS: Nash conditions

- (x^*, y^*) is a NE

$$\begin{aligned}(x^*)^T R y^* &\geq x^T R y^*, \text{ for } x \in X, \\ (x^*)^T C y^* &\geq (x^*)^T C y, \text{ for } y \in Y.\end{aligned}$$

- As $C = -R$ the second equation becomes

$$(x^*)^T R y^* \leq (x^*)^T R y, \text{ for } y \in Y.$$

- Combining both,

$$x^T R y^* \leq (x^*)^T R y^* \leq (x^*)^T R y, \text{ for } x \in X, y \in Y.$$

ZS: Nash conditions

- (x^*, y^*) is a NE


$$\begin{aligned}(x^*)^T R y^* &\geq x^T R y^*, \text{ for } x \in X, \\ (x^*)^T C y^* &\geq (x^*)^T C y, \text{ for } y \in Y.\end{aligned}$$

- As $C = -R$ the second equation becomes

$$(x^*)^T R y^* \leq (x^*)^T R y, \text{ for } y \in Y.$$

- Combining both,

$$x^T R y^* \leq (x^*)^T R y^* \leq (x^*)^T R y, \text{ for } x \in X, y \in Y.$$



i.e., (x^*, y^*) is a **saddle point**
of the function $x^T R y$ defined over $X \times Y$.

Minimax inequality

Minimax inequality

Theorem

For any function $\Phi : X \times Y \rightarrow \mathbb{R}$, we have

$$\sup_{x \in X} \inf_{y \in Y} \Phi(x, y) \leq \inf_{y \in Y} \sup_{x \in X} \Phi(x, y).$$

Minimax inequality

Theorem

For any function $\Phi : X \times Y \rightarrow \mathbb{R}$, we have

$$\sup_{x \in X} \inf_{y \in Y} \Phi(x, y) \leq \inf_{y \in Y} \sup_{x \in X} \Phi(x, y).$$

Proof.

Minimax inequality

Theorem

For any function $\Phi : X \times Y \rightarrow \mathbb{R}$, we have

$$\sup_{x \in X} \inf_{y \in Y} \Phi(x, y) \leq \inf_{y \in Y} \sup_{x \in X} \Phi(x, y).$$

Proof.

For every $x' \in X$, $\Phi(x', y) \leq \sup_{x \in X} \Phi(x, y)$

Minimax inequality

Theorem

For any function $\Phi : X \times Y \rightarrow \mathbb{R}$, we have

$$\sup_{x \in X} \inf_{y \in Y} \Phi(x, y) \leq \inf_{y \in Y} \sup_{x \in X} \Phi(x, y).$$

Proof.

For every $x' \in X$, $\Phi(x', y) \leq \sup_{x \in X} \Phi(x, y)$

$$\inf_{y \in Y} \Phi(x', y) \leq \inf_{y \in Y} \sup_{x \in X} \Phi(x, y).$$

Minimax inequality

Theorem

For any function $\Phi : X \times Y \rightarrow \mathbb{R}$, we have

$$\sup_{x \in X} \inf_{y \in Y} \Phi(x, y) \leq \inf_{y \in Y} \sup_{x \in X} \Phi(x, y).$$

Proof.

For every $x' \in X$, $\Phi(x', y) \leq \sup_{x \in X} \Phi(x, y)$

$$\inf_{y \in Y} \Phi(x', y) \leq \inf_{y \in Y} \sup_{x \in X} \Phi(x, y).$$

Taking the supremum over $x' \in X$ on the left hand-side we get the inequality. □

ZS: Nash conditions

We have

ZS: Nash conditions

We have

- $x^T R y^* \leq (x^*)^T R y^* \leq (x^*)^T R y$, for $x \in X$, $y \in Y$.

ZS: Nash conditions

We have

- $x^T R y^* \leq (x^*)^T R y^* \leq (x^*)^T R y$, for $x \in X$, $y \in Y$.
- Thus

$$\sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y$$

ZS: Nash conditions

We have

- $x^T R y^* \leq (x^*)^T R y^* \leq (x^*)^T R y$, for $x \in X$, $y \in Y$.
- Thus

$$\sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y$$

$$\inf_{y \in Y} \sup_{x \in X} x^T R y \leq \sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y \leq \sup_{x \in X} \inf_{y \in Y} x^T R y$$

ZS: Nash conditions

We have

- $x^T R y^* \leq (x^*)^T R y^* \leq (x^*)^T R y$, for $x \in X$, $y \in Y$.
- Thus

$$\sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y$$

$$\inf_{y \in Y} \sup_{x \in X} x^T R y \leq \sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y \leq \sup_{x \in X} \inf_{y \in Y} x^T R y$$

- Using the minimax inequality, we get

ZS: Nash conditions

We have

- $x^T R y^* \leq (x^*)^T R y^* \leq (x^*)^T R y$, for $x \in X$, $y \in Y$.
- Thus

$$\sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y$$

$$\inf_{y \in Y} \sup_{x \in X} x^T R y \leq \sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y \leq \sup_{x \in X} \inf_{y \in Y} x^T R y$$

- Using the minimax inequality, we get

$$\inf_{y \in Y} \sup_{x \in X} x^T R y = (x^*)^T R y^* = \sup_{x \in X} \inf_{y \in Y} x^T R y$$

ZS: Nash conditions

We have

- $x^T R y^* \leq (x^*)^T R y^* \leq (x^*)^T R y$, for $x \in X$, $y \in Y$.
- Thus

$$\sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y$$

$$\inf_{y \in Y} \sup_{x \in X} x^T R y \leq \sup_{x \in X} x^T R y^* \leq (x^*)^T R y^* \leq \inf_{y \in Y} (x^*)^T R y \leq \sup_{x \in X} \inf_{y \in Y} x^T R y$$

- Using the minimax inequality, we get

$$\inf_{y \in Y} \sup_{x \in X} x^T R y = (x^*)^T R y^* = \sup_{x \in X} \inf_{y \in Y} x^T R y$$

We refer to $\inf_{y \in Y} \sup_{x \in X} x^T R y$ as the **value of the game**.

ZS: algorithm for finding NE

- For a fixed y , we have

ZS: algorithm for finding NE

- For a fixed y , we have

$$\max_{x \in X} x^T R y = \max_{i=1, \dots, n} \{[Ry]_i\},$$

ZS: algorithm for finding NE

- For a fixed y , we have

$$\max_{x \in X} x^T R y = \max_{i=1, \dots, n} \{[Ry]_i\},$$

therefore

$$\min_{y \in Y} \max_{x \in X} x^T R y = \min_{y \in Y} \max\{[Ry]_1, \dots, [Ry]_n\}$$

ZS: algorithm for finding NE

- For a fixed y , we have

$$\max_{x \in X} x^T R y = \max_{i=1, \dots, n} \{[Ry]_i\},$$

therefore

$$\min_{y \in Y} \max_{x \in X} x^T R y = \min_{y \in Y} \max\{[Ry]_1, \dots, [Ry]_n\}$$

- So, both [the value of the game](#) and [a Nash equilibrium strategy for player 2](#) can be obtained by solving the linear programming problem:

ZS: algorithm for finding NE

- For a fixed y , we have

$$\max_{x \in X} x^T R y = \max_{i=1, \dots, n} \{[Ry]_i\},$$

therefore

$$\min_{y \in Y} \max_{x \in X} x^T R y = \min_{y \in Y} \max\{[Ry]_1, \dots, [Ry]_n\}$$

- So, both [the value of the game](#) and [a Nash equilibrium strategy for player 2](#) can be obtained by solving the linear programming problem:

$$\begin{aligned} \min v \\ v \mathbf{1}_n \geq R y, y \in Y. \end{aligned}$$

ZS: algorithm for finding NE

- Similarly, we have

$$\max_{x \in X} \min_{y \in Y} x^T R y = \max_{x \in X} \min \{ [R^T x]_1, \dots, [R^T x]_n \}$$

ZS: algorithm for finding NE

- Similarly, we have

$$\max_{x \in X} \min_{y \in Y} x^T R y = \max_{x \in X} \min \{ [R^T x]_1, \dots, [R^T x]_n \}$$

- So, a **Nash equilibrium strategy for player 1** can be obtained by solving the linear programming problem:

ZS: algorithm for finding NE

- Similarly, we have

$$\max_{x \in X} \min_{y \in Y} x^T R y = \max_{x \in X} \min \{ [R^T x]_1, \dots, [R^T x]_n \}$$

- So, a **Nash equilibrium strategy for player 1** can be obtained by solving the linear programming problem:

$$\begin{aligned} & \max w \\ & w \mathbf{1}_m \leq R^T x, x \in X. \end{aligned}$$

ZS: algorithm for finding NE

- Similarly, we have

$$\max_{x \in X} \min_{y \in Y} x^T R y = \max_{x \in X} \min \{ [R^T x]_1, \dots, [R^T x]_n \}$$

- So, a **Nash equilibrium strategy for player 1** can be obtained by solving the linear programming problem:

$$\begin{aligned} & \max w \\ & w \mathbf{1}_m \leq R^T x, x \in X. \end{aligned}$$

- LP can be solved efficiently, thus there is a **polynomial time algorithm for computing NE for zero-sum games**.

- 1 Linear Algebra formulation
- 2 Zero-sum games
- 3 The complexity of finding a NE**
- 4 An exact algorithm to compute NE
- 5 Other algorithms

PPAD

(Papadimitriou 94)

Polynomial Parity Argument on Directed Graphs

PPAD

(Papadimitriou 94)

Polynomial Parity Argument on Directed Graphs

- The class of all problems with guaranteed solution by use of the following graph-theoretic lemma

PPAD

(Papadimitriou 94)

Polynomial Parity Argument on Directed Graphs

- The class of all problems with guaranteed solution by use of the following graph-theoretic lemma

A directed graph with an unbalanced node (node with in-degree \neq outdegree) must have another.

PPAD

(Papadimitriou 94)

Polynomial Parity Argument on Directed Graphs

- The class of all problems with guaranteed solution by use of the following graph-theoretic lemma

A directed graph with an unbalanced node (node with in-degree \neq outdegree) must have another.

- Such problems are defined by an implicitly defined directed graph G and an unbalanced node u of G

PPAD

(Papadimitriou 94)

Polynomial Parity Argument on Directed Graphs

- The class of all problems with guaranteed solution by use of the following graph-theoretic lemma

A directed graph with an unbalanced node (node with in-degree \neq outdegree) must have another.

- Such problems are defined by an implicitly defined directed graph G and an unbalanced node u of G and the objective is finding another unbalanced node.

PPAD

(Papadimitriou 94)

Polynomial Parity Argument on Directed Graphs

- The class of all problems with guaranteed solution by use of the following graph-theoretic lemma
A directed graph with an unbalanced node (node with in-degree \neq outdegree) must have another.
- Such problems are defined by an implicitly defined directed graph G and an unbalanced node u of G and the objective is finding another unbalanced node.
- Usually G is huge but implicitly defined as the graphs defining solutions in local search algorithms.

PPAD

- The class PPAD contains interesting computational problems not known to be in P

PPAD

- The class PPAD contains interesting computational problems not known to be in P
has complete problems.

PPAD

- The class PPAD contains interesting computational problems not known to be in P has complete problems.
- But not a clear complexity cut.

A PPAD-complete problem

End-of-Line

A PPAD-complete problem

End-of-Line

Given an implicit representation of a graph G with vertices of degree at most 2 and a vertex $v \in G$, where v has in degree 0. Find a node $v' \neq v$, such that v' has out degree 0.

A PPAD-complete problem

End-of-Line

Given an implicit representation of a graph G with vertices of degree at most 2 and a vertex $v \in G$, where v has in degree 0. Find a node $v' \neq v$, such that v' has out degree 0.

- Since every node has degree 2, it is a collection of paths and cycles.

A PPAD-complete problem

End-of-Line

Given an implicit representation of a graph G with vertices of degree at most 2 and a vertex $v \in G$, where v has in degree 0. Find a node $v' \neq v$, such that v' has out degree 0.

- Since every node has degree 2, it is a collection of paths and cycles.
- We know that
Every directed graph with in/outdegree 1 and a source, has a sink.

A PPAD-complete problem

End-of-Line

Given an implicit representation of a graph G with vertices of degree at most 2 and a vertex $v \in G$, where v has in degree 0. Find a node $v' \neq v$, such that v' has out degree 0.

- Since every node has degree 2, it is a collection of paths and cycles.
- We know that
Every directed graph with in/outdegree 1 and a source, has a sink.
- Which guarantees that
the End-of-Line problem has always a solution.

End-of-Line: graph representation

- G is given implicitly by a circuit C
- C provides a predecessor and successor pair for each given vertex in G , i.e. $C(u) = (v, w)$.
- A special label indicates that a node has no predecessor/successor.

The complexity of finding a NE

Theorem (Daskalakis, Goldberg, Papadimitriou '06)

Finding a Nash equilibrium is PPAD-complete

The complexity of finding a NE

Theorem (Daskalakis, Goldberg, Papadimitriou '06)

Finding a Nash equilibrium is PPAD-complete

Theorem (Chen, Deng '06)

Finding a Nash equilibrium is PPAD-complete even in 2-player games.

The complexity of finding a NE

Theorem (Daskalakis, Goldberg, Papadimitriou '06)

Finding a Nash equilibrium is PPAD-complete

Theorem (Chen, Deng '06)

Finding a Nash equilibrium is PPAD-complete even in 2-player games.

- C. Daskalakis, P-W. Goldberg, C.H. Papadimitriou: [The complexity of computing a Nash equilibrium](#). SIAM J. Comput. 39(1): 195-259 (2009) first version STOC 2006
- X. Chen, X. Deng, S-H. Teng: [Settling the complexity of computing two-player Nash equilibria](#). J. ACM 56(3) (2009) first version FOCS 2006

- 1 Linear Algebra formulation
- 2 Zero-sum games
- 3 The complexity of finding a NE
- 4 An exact algorithm to compute NE**
- 5 Other algorithms

NE characterization

Theorem

In a strategic game in which each player has finitely many actions a mixed strategy profile σ^ is a NE iff, for each player i ,*

- the expected payoff, given σ_{-i} , to every action in the support of σ_i^* is the same*
- the expected payoff, given σ_{-i} , to every action not in the support of σ_i^* is at most the expected payoff on an action in the support of σ_i^* .*

NE conditions given support

Let $A \subseteq \{1, \dots, n\}$ and $B \subseteq \{1, \dots, m\}$.

The conditions for having a NE on this particular support can be written as follows:

$$\max \lambda_1 + \lambda_2$$

Subject to:

$$[R y]_i = \lambda_1, \text{ for } i \in A$$

$$[R y]_i \leq \lambda_1, \text{ for } i \notin A$$

$${}_j[C x] = \lambda_2, \text{ for } j \in B$$

$${}_j[C x] \leq \lambda_2, \text{ for } j \notin B$$

Iterating over all supports

- For every possible combination of supports $A \subseteq \{1, \dots, n\}$ and $B \subseteq \{1, \dots, m\}$.
Solve the set of simultaneous equations using linear programming.

Iterating over all supports

- For every possible combination of supports $A \subseteq \{1, \dots, n\}$ and $B \subseteq \{1, \dots, m\}$.
Solve the set of simultaneous equations using linear programming.
- This is an exact exponential time algorithm as the number of supports can be exponential.

Iterating over all supports

- For every possible combination of supports $A \subseteq \{1, \dots, n\}$ and $B \subseteq \{1, \dots, m\}$.
Solve the set of simultaneous equations using linear programming.
- This is an exact exponential time algorithm as the number of supports can be exponential.
- The same algorithm can be applied to a multiplayer game.
We would be able to compute a NE on rationals if such a NE exists.

- 1 Linear Algebra formulation
- 2 Zero-sum games
- 3 The complexity of finding a NE
- 4 An exact algorithm to compute NE
- 5 Other algorithms**

Other algorithms

- Lemke-Howson (1964) algorithm defines a polytope based on best response conditions and membership to the support and uses ideas similar to Simplex with a ad-hoc pivoting rule. (See slides by Ethan Kim)
- Lemke-Howson requires exponential time [[R. Savani, B. von Stengel, 2004](#)]).