

Efficiency of Nash Equilibria

Maria Serna

Fall 2019

- 1 Price of Anarchy/Stability
- 2 Selfish routing
- 3 Load Balancing game
- 4 Congestion games and variants
- 5 Affine Congestion games
- 6 References

Efficiency at equilibrium

- We have analyzed existence of PNE and NE
- The players' goals can be different from those of the society.
- Fixing a **social goal** an optimal situation is possible.
- How good/bad are NE with respect to this goal?

Efficiency at equilibrium

- We have analyzed existence of PNE and NE
- The players' goals can be different from those of the society.
- Fixing a **social goal** an optimal situation is possible.
- How good/bad are NE with respect to this goal?
- How far are NE for optimal social goal?

Efficiency at equilibrium

- We have analyzed existence of PNE and NE
- The players' goals can be different from those of the society.
- Fixing a **social goal** an optimal situation is possible.
- How good/bad are NE with respect to this goal?
- **How far are NE for optimal social goal?**
- To perform such an analysis for strategic games we have first to define a **global** function to optimize, this function is usually called the **social cost** or **social utility**.

Efficiency at equilibrium

- We have analyzed existence of PNE and NE
- The players' goals can be different from those of the society.
- Fixing a **social goal** an optimal situation is possible.
- How good/bad are NE with respect to this goal?
- **How far are NE for optimal social goal?**
- To perform such an analysis for strategic games we have first to define a **global** function to optimize, this function is usually called the **social cost** or **social utility**.
- Society is interested in minimizing the social cost or maximizing the social utility.

Social cost

- Consider a n -player game $\Gamma = (A_1, \dots, A_n, u_1, \dots, u_n)$.
- Let $A = A_1 \times \dots \times A_n$.
- Let $PNE(\Gamma)$ be the set of PNE of Γ .
- Let $NE(\Gamma)$ be the set of NE of Γ .

Social cost

- Consider a n -player game $\Gamma = (A_1, \dots, A_n, u_1, \dots, u_n)$.
- Let $A = A_1 \times \dots \times A_n$.
- Let $PNE(\Gamma)$ be the set of PNE of Γ .
- Let $NE(\Gamma)$ be the set of NE of Γ .
- Let $C : A \rightarrow \mathbb{R}$ be a social cost function.

C can be extended to mixed strategy profiles by computing the average under the joint product distribution.

Usual social cost functions

Usual social cost functions

- **Utilitarian social cost** : $C(s) = \sum_{i \in N} u_i(s)$.

Usual social cost functions

- **Utilitarian social cost** : $C(s) = \sum_{i \in N} u_i(s)$.
- **Egalitarian social cost**: $C(s) = \max_{i \in N} u_i(s)$.

Usual social cost functions

- **Utilitarian social cost** : $C(s) = \sum_{i \in N} u_i(s)$.
- **Egalitarian social cost**: $C(s) = \max_{i \in N} u_i(s)$.
- Game specific cost/utility defined by the model motivating the game.

Price of Anarchy/Stability

Price of Anarchy/Stability

The **Price of anarchy** of Γ is defined as

$$PoA(\Gamma) = \frac{\max_{\sigma \in NE(\Gamma)} C(\sigma)}{\min_{s \in A} C(s)}.$$

Price of Anarchy/Stability

The **Price of anarchy** of Γ is defined as

$$PoA(\Gamma) = \frac{\max_{\sigma \in NE(\Gamma)} C(\sigma)}{\min_{s \in A} C(s)}.$$

The **Price of stability** of Γ is defined as

$$PoS(\Gamma) = \frac{\min_{\sigma \in NE(\Gamma)} C(\sigma)}{\min_{s \in A} C(s)}.$$

Price of Anarchy/Stability

The **Price of anarchy** of Γ is defined as

$$PoA(\Gamma) = \frac{\max_{\sigma \in NE(\Gamma)} C(\sigma)}{\min_{s \in A} C(s)}.$$

The **Price of stability** of Γ is defined as

$$PoS(\Gamma) = \frac{\min_{\sigma \in NE(\Gamma)} C(\sigma)}{\min_{s \in A} C(s)}.$$

For social utility functions the terms are inverted in the definition.

Price of Anarchy/Stability

Price of Anarchy/Stability

- For games having a PNE, we might be interested in those values over $PNE(\Gamma)$ instead of $NE(\Gamma)$.

Price of Anarchy/Stability

- For games having a PNE, we might be interested in those values over $PNE(\Gamma)$ instead of $NE(\Gamma)$.
- For families of games, we might be interested in analyzing PoA and PoS as a function of some parameter. For example the number of players.

Price of Anarchy/Stability


- For games having a PNE, we might be interested in those values over $PNE(\Gamma)$ instead of $NE(\Gamma)$.
- For families of games, we might be interested in analyzing PoA and PoS as a function of some parameter. For example the number of players.
- PoA measures the **worst decentralized** equilibrium scenario, the one giving the maximum system degradation.

Price of Anarchy/Stability

- For games having a PNE, we might be interested in those values over $PNE(\Gamma)$ instead of $NE(\Gamma)$.
- For families of games, we might be interested in analyzing PoA and PoS as a function of some parameter. For example the number of players.
- PoA measures the **worst decentralized** equilibrium scenario, the one giving the maximum system degradation.
- PoS measures the **best decentralized** equilibrium scenario, the one giving the best possible degradation.


Network congestion game

- 4000 drivers drive from A to B on



Network congestion game


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.

Network congestion game


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.
- Optimal social cost is reached when half of the drivers take $A - U - B$ and the other half $A - R - B$ with social cost 65.

Network congestion game


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.
- Optimal social cost is reached when half of the drivers take $A - U - B$ and the other half $A - R - B$ with social cost 65.
- In the NE

Network congestion game


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.
- Optimal social cost is reached when half of the drivers take $A - U - B$ and the other half $A - R - B$ with social cost 65.
- In the NE half of the drivers take $A - U - B$ and the other half $A - R - B$.

Network congestion game


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.
- Optimal social cost is reached when half of the drivers take $A - U - B$ and the other half $A - R - B$ with social cost 65.
- In the NE half of the drivers take $A - U - B$ and the other half $A - R - B$.
- $PoA = PoS = 65/65 = 1$


Braess' Network

- 4000 drivers drive from A to B on



Braess' Network


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.

Braess' Network


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.
- Optimal social cost is reached when half of the drivers take $A - U - B$ and the other half $A - R - B$ with social cost 65.

Braess' Network


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.
- Optimal social cost is reached when half of the drivers take $A - U - B$ and the other half $A - R - B$ with social cost 65.
- In the NE

Braess' Network


- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.
- Optimal social cost is reached when half of the drivers take $A - U - B$ and the other half $A - R - B$ with social cost 65.
- In the NE all drivers take $A - U - R - B$ with social cost 80.

Braess' Network

- 4000 drivers drive from A to B on



- Set the social cost to be the **maximum travel time**.
- Optimal social cost is reached when half of the drivers take $A - U - B$ and the other half $A - R - B$ with social cost 65.
- In the NE all drivers take $A - U - R - B$ with social cost 80.
- $PoA = PoS = 80/65 = 16/13$


- 1 Price of Anarchy/Stability
- 2 Selfish routing**
- 3 Load Balancing game
- 4 Congestion games and variants
- 5 Affine Congestion games
- 6 References

Selfish Routing: an Example

- Total traffic is $r = 1$.
- Network (with delay functions on arcs)


Selfish Routing: an Example

- Total traffic is $r = 1$.
- Network (with delay functions on arcs)



Selfish Routing: an Example


- Total traffic is $r = 1$.
- Network (with delay functions on arcs)



- Player's objective going from $s = A$ to $s = B$ with minimum delay.

Selfish Routing: an Example

- Total traffic is $r = 1$.
- Network (with delay functions on arcs)



- Player's objective going from $s = A$ to $s = B$ with minimum delay.
- Strategy profiles: flows from A to B with total flow $r = 1$

Selfish routing: strategy profiles

Selfish routing: strategy profiles

- Traffic as Flows:

Selfish routing: strategy profiles


- Traffic as Flows:
 - A **flow** f giving the routing of traffic.

Selfish routing: strategy profiles

- Traffic as Flows:
 - A **flow** f giving the routing of traffic.
 - Recall that a flow must preserve flow in = flow out except for sources/sinks.


Selfish routing: strategy profiles

- Traffic as Flows:
 - A **flow** f giving the routing of traffic.
 - Recall that a flow must preserve flow in = flow out except for sources/sinks.



Selfish routing: strategy profiles

- Traffic as Flows:
 - A **flow** f giving the routing of traffic.
 - Recall that a flow must preserve flow in = flow out except for sources/sinks.



- Notation: for a path P and a feasible flow f , $C^P(f)$ denotes the cost corresponding to the traffic routed through P by f .

Selfish routing: equilibria


Theorem

A flow is a Nash equilibrium (or is a Nash flow) if all flow is routed on min-latency paths (given current edge congestion)

Selfish routing: equilibria

Theorem


A flow is a Nash equilibrium (or is a Nash flow) if all flow is routed on min-latency paths (given current edge congestion)



Selfish routing: equilibria

Theorem

A flow is a Nash equilibrium (or is a Nash flow) if all flow is routed on min-latency paths (given current edge congestion)



For f to be a Nash equilibrium, all $A - B$ paths should have minimum latency, so $f = 1/2$.

Selfish routing: equilibria

Selfish routing: equilibria

Theorem

A feasible flow x is an equilibrium flow iff for any feasible flow y

$$\sum_{e \in E} d_e(x[e])x[e] \leq \sum_{e \in E} d_e(x[e])y[e].$$

Selfish routing: equilibria

Theorem

A feasible flow x is an equilibrium flow iff for any feasible flow y

$$\sum_{e \in E} d_e(x[e])x[e] \leq \sum_{e \in E} d_e(x[e])y[e].$$

- Called **Variational Inequality** (Smith 79 and Dafermos 80)

Selfish routing: equilibria

Theorem

A feasible flow x is an equilibrium flow iff for any feasible flow y

$$\sum_{e \in E} d_e(x[e])x[e] \leq \sum_{e \in E} d_e(x[e])y[e].$$

- Called **Variational Inequality** (Smith 79 and Dafermos 80)
- As a consequence all Nash flows have the same cost per edge.

Selfish routing: equilibria

Theorem

A feasible flow x is an equilibrium flow iff for any feasible flow y

$$\sum_{e \in E} d_e(x[e])x[e] \leq \sum_{e \in E} d_e(x[e])y[e].$$

- Called **Variational Inequality** (Smith 79 and Dafermos 80)
- As a consequence all Nash flows have the same cost per edge.
- Do PNE exist?

Selfish routing: equilibria existence

- As for the atomic case we can consider a potential, for a given flow x

$$\Psi(x) = \sum_{e \in E} \int_0^{x[e]} d_e(u) du$$

Selfish routing: equilibria existence

- As for the atomic case we can consider a potential, for a given flow x

$$\Psi(x) = \sum_{e \in E} \int_0^{x[e]} d_e(u) du$$

Theorem

A feasible flow x is an equilibrium flow iff x is a minimum of Ψ over the set of feasible flows.

Selfish routing

- Social cost: **maximum travel time** egalitarian
 - By the characterization of Nash flows all NE have the same cost.
 - $PoA = PoS = \text{cost NE} / \text{opt}$

Selfish routing

- Social cost: **maximum travel time** egalitarian
 - By the characterization of Nash flows all NE have the same cost.
 - $PoA = PoS = \text{cost NE} / \text{opt}$
- Other social cost? utilitarian


A natural one is the **total travel time**

Selfish routing: total routing time

The **cost** $C(f)$ of flow f is the sum of all delays incurred by traffic.

Selfish routing: total routing time


The **cost** $C(f)$ of flow f is the sum of all delays incurred by traffic.



$$C(f) = f(f + 1) + (1 - f)2.$$

Selfish routing: total routing time

The **cost** $C(f)$ of flow f is the sum of all delays incurred by traffic.



$$C(f) = f(f + 1) + (1 - f)2.$$

Formally, if $d_P(f)$ is the sum of latencies of edges in a path P :


$$C(f) = \sum_P f_P d_P(f)$$

Selfish routing: inefficiency of Nash flows

- Nash flows do not minimize total latency


Selfish routing: inefficiency of Nash flows

- Nash flows do not minimize total latency



Selfish routing: inefficiency of Nash flows


- Nash flows do not minimize total latency



- Cost: $f^2 + 1 - f + 0.5$

Selfish routing: inefficiency of Nash flows


- Nash flows do not minimize total latency



- Cost: $f^2 + 1 - f + 0.5$
- Optimal cost $0.25 + 0.5 + 0.5 = 1.25$

Selfish routing: inefficiency of Nash flows

- Nash flows do not minimize total latency




- Cost: $f^2 + 1 - f + 0.5$
- Optimal cost $0.25 + 0.5 + 0.5 = 1.25$
- Nash flow has cost 1.5

Selfish routing: inefficiency of Nash flows

- An extreme case:


Selfish routing: inefficiency of Nash flows

- An extreme case:



Selfish routing: inefficiency of Nash flows


- An extreme case:



- Cost: $(1 - f)^d + f$

Selfish routing: inefficiency of Nash flows


- An extreme case:



- Cost: $(1 - f)^d + f$
- For d large enough, as $r = 1$, optimal cost $\epsilon + (1 - \epsilon)^d$ where $\epsilon \rightarrow 0$ as $d \rightarrow \infty$

Selfish routing: inefficiency of Nash flows


- An extreme case:



- Cost: $(1 - f)^d + f$
- For d large enough, as $r = 1$, optimal cost $\epsilon + (1 - \epsilon)^d$ where $\epsilon \rightarrow 0$ as $d \rightarrow \infty$
- Nash flow has cost 1

Selfish routing: inefficiency of Nash flows

- An extreme case:



- Cost: $(1 - f)^d + f$
- For d large enough, as $r = 1$, optimal cost $\epsilon + (1 - \epsilon)^d$ where $\epsilon \rightarrow 0$ as $d \rightarrow \infty$
- Nash flow has cost 1
- Unbounded PoA: Nash flow can cost arbitrarily more than the optimal (mincost) flow even if latency functions are polynomials

- 1 Price of Anarchy/Stability
- 2 Selfish routing
- 3 Load Balancing game**
- 4 Congestion games and variants
- 5 Affine Congestion games
- 6 References

Load Balancing game

- There are m servers and n jobs. Job i has load p_i .
- The game has n players, corresponding to the n jobs.
- Each player has to decide the server that will process its job.
 $A_i = \{1, \dots, m\}$
- The response time of server j is proportional to its load

$$L_j(s) = \sum_{i|s_i=j} p_i.$$

- Each job wants to be assigned to the server that minimizes its response time:

$$c_i(s) = L_{s_i}(s).$$

Load Balancing game: PNE?

Consider the best response dynamic

- Start with an arbitrary state.
- A node (or several) chooses a best strategy, one that maximizes its own payoff, given the current choices of the others

Load Balancing game: PNE?

Consider the best response dynamic

- Start with an arbitrary state.
- A node (or several) chooses a best strategy, one that maximizes its own payoff, given the current choices of the others
- How to prove that such a process converges to a PNE?

Load Balancing game: PNE?

Consider the best response dynamic

- Start with an arbitrary state.
- A node (or several) chooses a best strategy, one that maximizes its own payoff, given the current choices of the others
- How to prove that such a process converges to a PNE?
- Seek for an adequate kind of **potential** function.

Load Balancing game: PNE?

BR-inspired-algorithm

- Order the servers with decreasing load (i.e., the decreasing response time):

$$L_1 \geq L_2 \geq \dots \geq L_m.$$

- Job i moves from server j to k , $L_k + p_i < L_j$.
- We must have $L_1 \geq \dots \geq L_j \geq \dots \geq L_k \geq \dots \geq L_m$.
- Thus, $L_j - p_i, L_k + p_i < L_j$

Load Balancing game: PNE?

BR-inspired-algorithm

- Order the servers with decreasing load (i.e., the decreasing response time):

$$L_1 \geq L_2 \geq \dots \geq L_m.$$

- Job i moves from server j to k , $L_k + p_i < L_j$.
- We must have $L_1 \geq \dots \geq L_j \geq \dots \geq L_k \geq \dots \geq L_m$.
- Thus, $L_j - p_i, L_k + p_i < L_j$
- Reorder the servers by decreasing load and repeat the process until no job can move.

Load Balancing game: PNE?

- Does the algorithm converge?

Load Balancing game: PNE?

- Does the algorithm converge?
- There are a finite number of (possibly exponential) assignments of jobs to servers.

Load Balancing game: PNE?

- Does the algorithm converge?
- There are a finite number of (possibly exponential) assignments of jobs to servers.
- At each step the sorted load sequence **decreases lexicographically!**

Load Balancing game: PNE?

- Does the algorithm converge?
- There are a finite number of (possibly exponential) assignments of jobs to servers.
- At each step the sorted load sequence **decreases lexicographically!**
- So BR-inspired-algorithm terminates (although it can be rather slow).

Load Balancing game: PNE?

- Does the algorithm converge?
- There are a finite number of (possibly exponential) assignments of jobs to servers.
- At each step the sorted load sequence **decreases lexicographically!**
- So BR-inspired-algorithm terminates (although it can be rather slow).
- The load balancing game has a PNE.

Load Balancing game: Social cost

- The natural social cost is the **total finish time** i.e., the maximum of the server's loads

$$c(s) = \max_{j=1}^m L_j.$$

- How bad/good is a PNE?

Load Balancing game: PoS

- Let s be an assignment with optimal cost.
- Is s a PNE?

Load Balancing game: PoS

- Let s be an assignment with optimal cost.
- Is s a PNE?
- not necessarily, no player in the worst server can improve, however other players can get a better benefit.

Load Balancing game: PoS

- Let s be an assignment with optimal cost.
- Is s a PNE?
- not necessarily, no player in the worst server can improve, however other players can get a better benefit.
- However, starting from an optimal solution the BR-inspired-algorithm terminates on a PNE with the same maximum load.

Load Balancing game: PoS

- Let s be an assignment with optimal cost.
- Is s a PNE?
- not necessarily, no player in the worst server can improve, however other players can get a better benefit.
- However, starting from an optimal solution the BR-inspired-algorithm terminates on a PNE with the same maximum load.
- Therefore, $PoS(\Gamma) = 1$.

Load Balancing game: PoA

Theorem

The max load of a Pure Nash equilibrium s is within twice the max load of an optimum assignment, i.e.,

$$C(s) \leq 2 \min_{s'} C(s').$$

Which will give $PoA(\Gamma) \leq 2$.

Load Balancing game: PoA bound

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.
- In an opt solution s' , i is assigned to some server,

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.
- In an opt solution s' , i is assigned to some server, so $C(s') \geq p_i$.

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.
- In an opt solution s' , i is assigned to some server, so $C(s') \geq p_i$.
- $\sum_k L_k$ is the total processing time for an assignment.

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.
- In an opt solution s' , i is assigned to some server, so $C(s') \geq p_i$.
- $\sum_k L_k$ is the total processing time for an assignment.
The best possible algorithm is to evenly partition them among m servers (if possible), thus

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.
- In an opt solution s' , i is assigned to some server, so $C(s') \geq p_i$.
- $\sum_k L_k$ is the total processing time for an assignment.
The best possible algorithm is to evenly partition them among m servers (if possible), thus $C(s') \geq \sum_k L_k/m = (\sum_\ell p_\ell)/m$.

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.
- In an opt solution s' , i is assigned to some server, so $C(s') \geq p_i$.
- $\sum_k L_k$ is the total processing time for an assignment.
The best possible algorithm is to evenly partition them among m servers (if possible), thus $C(s') \geq \sum_k L_k/m = (\sum_\ell p_\ell)/m$.
- We get
$$C(s) = L_j \leq (\sum_k L_k)/m + p_i$$

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.
- In an opt solution s' , i is assigned to some server, so $C(s') \geq p_i$.
- $\sum_k L_k$ is the total processing time for an assignment.
 The best possible algorithm is to evenly partition them among m servers (if possible), thus $C(s') \geq \sum_k L_k/m = (\sum_\ell p_\ell)/m$.
- We get

$$C(s) = L_j \leq (\sum_k L_k)/m + p_i \leq (\sum_\ell p_\ell)/m + p_i$$

Load Balancing game: PoA bound

- Let s be a PNE
- Let i be a job assigned in s to the max loaded server j .
 - $L_j \leq L_k + p_i$, for all other server k .
 - Summing over all servers, we get $L_j \leq (\sum_k L_k)/m + p_i$.
- In an opt solution s' , i is assigned to some server, so $C(s') \geq p_i$.
- $\sum_k L_k$ is the total processing time for an assignment.
 The best possible algorithm is to evenly partition them among m servers (if possible), thus $C(s') \geq \sum_k L_k/m = (\sum_\ell p_\ell)/m$.
- We get

$$C(s) = L_j \leq (\sum_k L_k)/m + p_i \leq (\sum_\ell p_\ell)/m + p_i \leq C(s') + C(s').$$

- 1 Price of Anarchy/Stability
- 2 Selfish routing
- 3 Load Balancing game
- 4 Congestion games and variants**
- 5 Affine Congestion games
- 6 References

Congestion games

Congestion games

A **congestion game** $(E, N, (d_e)_{e \in E}, (c_i)_{i \in N})$

- is defined on a finite set E of resources and
- has n players
- using a delay function d_e mapping \mathbb{N} to the integers, for each resource e .
- The actions for each player are subsets of E .
- The cost functions are the following:

$$c_i(a_1, \dots, a_n) = \sum_{e \in a_i} d_e(f_e(a_1, \dots, a_n))$$

being $f_e(a_1, \dots, a_n) = |\{i \mid e \in a_i\}|$.

Weighted congestion games

Weighted congestion games

A **weighted congestion game** $(E, N, (d_e)_{e \in E}, (c_i)_{i \in N}, (w_i)_{i \in N})$

- is defined on a finite set E of resources and
- has n players. Player i has an associated **natural weight** w_i .
- Using a delay function d_e mapping \mathbb{N} to the integers, for each resource e .
- The actions for each player are subsets of E .
- The cost functions are the following:

$$c_i(a_1, \dots, a_n) = \sum_{e \in a_i} d_e(f_e(a_1, \dots, a_n))$$

being $f_e(a_1, \dots, a_n) = \sum_{\{i | e \in a_i\}} w_i$.

Network weighted congestion games

Network weighted congestion games

A **network weighted congestion game** is defined on a directed graph $G = (V, E)$, $(N, G, (d_e)_{e \in E}, (c_i)_{i \in N}, (w_i)_{i \in N}, (s_i)_{i \in N}, (t_i)_{i \in N})$.

- The resources are the arcs in G .
- The game has n players. Player i has an associated natural weight w_i .
- Using a delay function d_e mapping \mathbb{N} to the integers, for each arc $e \in E$.
- The action set for player i is the set of (s_i, t_i) -paths in G .
- The cost functions are the following:

$$c_i(a_1, \dots, a_n) = \sum_{e \in a_i} d_e(f_e(a_1, \dots, a_n))$$

being $f_e(a_1, \dots, a_n) = \sum_{\{i | e \in a_i\}} w_i$.

Another family: Fair Cost Sharing Games

Another family: Fair Cost Sharing Games

A **fair cost sharing game** $(E, N, (c_e)_{e \in E})$

- is defined on a finite set E of resources and
- has n players
- a **fixed cost** c_e , for each resource e .
- The actions for each player are subsets of E .
- The cost functions are the following:

$$c_i(a_1, \dots, a_n) = \sum_{e \in a_i} \frac{c_e}{f_e(a_1, \dots, a_n)}$$

being $f_e(a_1, \dots, a_n) = |\{i \mid e \in a_i\}|$.

Congestion games terminology

Congestion games terminology

- unweighted (vs. weighted)

Congestion games terminology

- unweighted (vs. weighted): $w_i = 1$.

Congestion games terminology

- unweighted (vs. weighted): $w_i = 1$.
- symmetric (vs. non-symmetric) strategies:

Congestion games terminology

- **unweighted (vs. weighted)**: $w_i = 1$.
- **symmetric (vs. non-symmetric) strategies**: all the players have the same set of actions.

Congestion games terminology

- unweighted (vs. weighted): $w_i = 1$.
- symmetric (vs. non-symmetric) strategies: all the players have the same set of actions.
- symmetric congestion games:

Congestion games terminology

- **unweighted (vs. weighted)**: $w_i = 1$.
- **symmetric (vs. non-symmetric) strategies**: all the players have the same set of actions.
- **symmetric congestion games**: unweighted with symmetric strategies.

Congestion games terminology

- **unweighted (vs. weighted)**: $w_i = 1$.
- **symmetric (vs. non-symmetric) strategies**: all the players have the same set of actions.
- **symmetric congestion games**: unweighted with symmetric strategies.
- **singleton congestion games**:

Congestion games terminology

- **unweighted (vs. weighted)**: $w_i = 1$.
- **symmetric (vs. non-symmetric) strategies**: all the players have the same set of actions.
- **symmetric congestion games**: unweighted with symmetric strategies.
- **singleton congestion games**: all possible actions have only one resource.

Congestion games terminology

- **unweighted (vs. weighted)**: $w_i = 1$.
- **symmetric (vs. non-symmetric) strategies**: all the players have the same set of actions.
- **symmetric congestion games**: unweighted with symmetric strategies.
- **singleton congestion games**: all possible actions have only one resource.
- **nonatomic network congestion games (vs. atomic)**

Congestion games terminology

- **unweighted (vs. weighted)**: $w_i = 1$.
- **symmetric (vs. non-symmetric) strategies**: all the players have the same set of actions.
- **symmetric congestion games**: unweighted with symmetric strategies.
- **singleton congestion games**: all possible actions have only one resource.
- **nonatomic network congestion games (vs. atomic)**
In **nonatomic** congestion games the number of players is infinite and each player controls an infinitesimal weight of the total traffic.

Congestion games terminology

- **unweighted (vs. weighted)**: $w_i = 1$.
- **symmetric (vs. non-symmetric) strategies**: all the players have the same set of actions.
- **symmetric congestion games**: unweighted with symmetric strategies.
- **singleton congestion games**: all possible actions have only one resource.
- **nonatomic network congestion games (vs. atomic)**
In **nonatomic** congestion games the number of players is infinite and each player controls an infinitesimal weight of the total traffic. Named also **Selfish routing games**.

PNE in Weighted Congestion Games


- There are weighted network congestion games without PNE

PNE in Weighted Congestion Games

- There are weighted network congestion games without PNE
- The following network with 2 players having weights $w_1 = 1$ and $w_2 = 2$

PNE in Weighted Congestion Games

- There are weighted network congestion games without PNE
- The following network with 2 players having weights $w_1 = 1$ and $w_2 = 2$



Not always PNE in Weighted Congestion Games

Not always PNE in Weighted Congestion Games

s_{-i}	BR_1	BR_2
$P_1 : s \rightarrow t$	P_4	P_2
$P_2 : s \rightarrow v \rightarrow t$	P_4	P_4
$P_3 : s \rightarrow w \rightarrow t$	P_1	P_2
$P_4 : s \rightarrow v \rightarrow w \rightarrow t$	P_1	P_3

Not always PNE in Weighted Congestion Games

s_{-i}	BR_1	BR_2
$P_1 : s \rightarrow t$	P_4	P_2
$P_2 : s \rightarrow v \rightarrow t$	P_4	P_4
$P_3 : s \rightarrow w \rightarrow t$	P_1	P_2
$P_4 : s \rightarrow v \rightarrow w \rightarrow t$	P_1	P_3

Therefore the game has no PNE

- 1 Price of Anarchy/Stability
- 2 Selfish routing
- 3 Load Balancing game
- 4 Congestion games and variants
- 5 Affine Congestion games**
- 6 References

Affine congestion games

Consider unweighted congestion games such that the delay functions are **affine** functions, i.e., for each resource e ,

$$d_e(x) = a_e x + b_e,$$

for some $a_e, b_e > 0$.

Affine congestion games

Consider unweighted congestion games such that the delay functions are **affine** functions, i.e., for each resource e ,

$$d_e(x) = a_e x + b_e,$$

for some $a_e, b_e > 0$.

Let C be the usual social cost:

$$C(s) = \sum_{e \in E} d_e(f_e(s))$$

PNE in Affine Congestion Games

PNE in Affine Congestion Games

- For affine delay functions PNE always exist

PNE in Affine Congestion Games

- For affine delay functions PNE always exist
Show that the following $\Phi(s)$ is a **weighted** potential function

PNE in Affine Congestion Games

- For affine delay functions PNE always exist
Show that the following $\Phi(s)$ is a **weighted** potential function

$$U(s) = \sum_{i \in N} w_i \sum_{e \in s_i} (a_e w_i + b_e) \quad C(s) = \sum_{i \in N} w_i c_i(s)$$

$$\Phi(s) = (C(s) + U(s))/2.$$

PNE in Affine Congestion Games

- For affine delay functions PNE always exist
Show that the following $\Phi(s)$ is a **weighted** potential function

$$U(s) = \sum_{i \in N} w_i \sum_{e \in s_i} (a_e w_i + b_e) \quad C(s) = \sum_{i \in N} w_i c_i(s)$$
$$\Phi(s) = (C(s) + U(s))/2.$$

You should be able to show that

$$\Phi(s') - \Phi(s) = w_i(c_i(s') - c_i(s)).$$

Smoothness

A game is called (λ, μ) -smooth, for $\lambda > 0$ and $\mu \leq 1$ if, for every pair of strategy profiles s and s' , we have

$$\sum_{i \in N} c_i(s_{-i}, s'_i) \leq \lambda C(s') + \mu C(s).$$

Smoothness

A game is called (λ, μ) -smooth, for $\lambda > 0$ and $\mu \leq 1$ if, for every pair of strategy profiles s and s' , we have

$$\sum_{i \in N} c_i(s_{-i}, s'_i) \leq \lambda C(s') + \mu C(s).$$

Smoothness directly gives a bound for the PoA:

Smoothness

A game is called (λ, μ) -smooth, for $\lambda > 0$ and $\mu \leq 1$ if, for every pair of strategy profiles s and s' , we have

$$\sum_{i \in N} c_i(s_{-i}, s'_i) \leq \lambda C(s') + \mu C(s).$$

Smoothness directly gives a bound for the PoA:

Theorem

In a (λ, μ) -smooth game, the PoA for PNE is at most $\frac{\lambda}{1-\mu}$.

Proof of smoothness bound on PoA

Let s be the worst PNE and s^* be an optimum solution.

$$\begin{aligned} C(s) &= \sum_{i \in N} c_i(s) \leq \sum_{i \in N} c_i(s_{-i}, s_i^*) \\ &\leq \lambda C(s^*) + \mu C(s) \end{aligned}$$

Subtracting $\mu C(s)$ on both sides gives

$$(1 - \mu)C(s) \leq \lambda C(s^*).$$

Theorem

Every congestion game with affine delay functions is $(5/3, 1/3)$ -smooth. Thus, $PoA \leq 5/2$.

Theorem

Every congestion game with affine delay functions is $(5/3, 1/3)$ -smooth. Thus, $PoA \leq 5/2$.

The proof uses a technical lemma:

Lemma (Christodoulou, Koutsoupias, 2005)

For all integers y, z we have

$$y(z + 1) \leq \frac{5}{3}y^2 + \frac{1}{3}z^2.$$

Proof of smoothness for affine functions

Recall that $d_e(x) = a_e x + b_e$. Note that using the Lemma

$$a_e y(z+1) + b_e y \leq a_e \left(\frac{5}{3} y^2 + \frac{1}{3} z^2 \right) + b_e y = \frac{5}{3} (a_e y^2 + b_e y) + \frac{1}{3} (a_e z^2 + b_e z).$$

Proof of smoothness for affine functions

Recall that $d_e(x) = a_e x + b_e$. Note that using the Lemma

$$a_e y(z+1) + b_e y \leq a_e \left(\frac{5}{3} y^2 + \frac{1}{3} z^2 \right) + b_e y = \frac{5}{3} (a_e y^2 + b_e y) + \frac{1}{3} (a_e z^2 + b_e z).$$

Taking $y = f_e(s^*)$ and $z = f_e(s)$ we get

$$(a_e(f_e(s)+1) + b_e)f_e(s^*) \leq \frac{5}{3} (a_e f_e(s^*) + b_e)f_e(s^*) + \frac{1}{3} (a_e f_e(s) + b_e)f_e(s).$$

Proof of smoothness for affine functions

Recall that $d_e(x) = a_e x + b_e$. Note that using the Lemma

$$a_e y(z+1) + b_e y \leq a_e \left(\frac{5}{3} y^2 + \frac{1}{3} z^2 \right) + b_e y = \frac{5}{3} (a_e y^2 + b_e y) + \frac{1}{3} (a_e z^2 + b_e z).$$

Taking $y = f_e(s^*)$ and $z = f_e(s)$ we get

$$(a_e(f_e(s)+1) + b_e)f_e(s^*) \leq \frac{5}{3} (a_e f_e(s^*) + b_e)f_e(s^*) + \frac{1}{3} (a_e f_e(s) + b_e)f_e(s).$$

Summing up all the inequalities

$$\sum_{e \in E} (a_e(f_e(s) + 1) + b_e)f_e(s^*) \leq \frac{5}{3} C(s^*) + \frac{1}{3} C(s).$$

Proof of smoothness for affine functions

$$\sum_{e \in E} (a_e(f_e(s) + 1) + b_e)f_e(s^*) \leq \frac{5}{3}C(s^*) + \frac{1}{3}C(s).$$

Proof of smoothness for affine functions

$$\sum_{e \in E} (a_e(f_e(s) + 1) + b_e)f_e(s^*) \leq \frac{5}{3}C(s^*) + \frac{1}{3}C(s).$$

But,

$$\sum_{i \in N} c_i(s_{-i}, s_i^*) \leq \sum_{e \in E} (a_e(f_e(s) + 1) + b_e)f_e(s^*)$$

as there are at most $f_e(s^*)$ players that might move to resource r . Each of them by unilaterally deviating incur a delay of $(a_e(f_e(s) + 1) + b_e)$.

Proof of smoothness for affine functions

$$\sum_{e \in E} (a_e(f_e(s) + 1) + b_e)f_e(s^*) \leq \frac{5}{3}C(s^*) + \frac{1}{3}C(s).$$

But,

$$\sum_{i \in N} c_i(s_{-i}, s_i^*) \leq \sum_{e \in E} (a_e(f_e(s) + 1) + b_e)f_e(s^*)$$

as there are at most $f_e(s^*)$ players that might move to resource r .
 Each of them by unilaterally deviating incur a delay of $(a_e(f_e(s) + 1) + b_e)$.

This gives the $(5/3, 1/3)$ -smoothness.

- 1 Price of Anarchy/Stability
- 2 Selfish routing
- 3 Load Balancing game
- 4 Congestion games and variants
- 5 Affine Congestion games
- 6 References

References

- Chapters 18 and 19.3 in the AGT book. (PoA and PoS bounds).
- B. Awerbuch, Y. Azar, A. Epstein. The Price of Routing Unsplittable Flow. STOC 2005. (PoA for pure NE in congestion games).
- G. Christodoulou, E. Koutsoupias. The Price of Anarchy of finite Congestion Games. STOC 2005. (PoA for pure NE in congestion games)
- T. Roughgarden. Intrinsic Robustness of the Price of Anarchy. STOC 2009. (Smoothness Framework and Unification of Previous Results)
- D. Fotakis. A Selective Tour Through Congestion Games, LNCS 2015.